

Think you have Solved Commonsense Reasoning? Try HellaSwagUltra

Wenhan Han¹ Yifan Zhang² Binbin Liu³

Mykola Pechenizkiy¹ Yin Zheng³ Meng Fang^{4,1}

¹Eindhoven University of Technology ²Nanjing University ³Tsinghua University

⁴University of Liverpool

Correspondence: meng.fang@liverpool.ac.uk, yzheng3xg@gmail.com

Abstract

As LLMs advance, popular commonsense and NLU benchmarks have become saturated, while multilingual evaluation still lags. Existing benchmarks cover few languages, and translation-based multilingual benchmarks can introduce bias. We introduce HellaSwagUltra, a multilingual commonsense and NLU benchmark with substantial local cultural knowledge, built by an automated and continuously extensible pipeline. Rather than explicitly prompting for reasoning, each example embeds two commonsense or local knowledge facts in the context, and answer options include subtle cues about whether those facts are violated, requiring models to choose the most plausible continuation. We release a 2,926-item HellaSwagUltra-Gold subset fully reviewed and corrected by language experts, with ongoing updates. On this human-verified subset, even the strongest proprietary model reaches only 62.5 percent accuracy, while GPT 4o and leading open source models stay around 40 to 50 percent. These results show multilingual commonsense reasoning remains open, and we release the dataset and pipeline to support future work. Our dataset and pipeline are publicly available at <https://hellaswagultra.github.io/>.

1 Introduction

LLM development is shifting decisively toward multilinguality. Modern academic and commercial models increasingly claim competence across dozens or even hundreds of languages, moving beyond an English-centric paradigm. For example, Gemma3 (Team et al., 2025) reports coverage of 140 plus languages, and Qwen3 (Yang et al., 2025) supports 119 languages and dialects. Closed-source systems such as ChatGPT (OpenAI et al., 2024), Claude ¹, and Gemini (Team et al., 2024a)

also emphasize multilingual performance, although exact coverage is often undisclosed.

Multilingual evaluation has not kept pace with rapid LLM progress, and multilingual commonsense reasoning remains especially underexplored. The core obstacle is data scarcity: for most languages beyond English, both data quality and scale are far weaker. A common workaround is to translate English benchmarks (Lai et al., 2023; Huang et al., 2025; Singh et al., 2025), but this often introduces translation errors and cultural bias. Recent native multilingual suites such as CMMLU (Li et al., 2024), INCLUDE (Romanou et al., 2024), and MultiLoKo (Hupkes and Bogoychev, 2025) help reduce these issues, yet they are largely sourced from Wikipedia style text or exam questions. As a result, they emphasize factual knowledge and still fall short on the key capabilities that multilingual evaluation should stress: natural language understanding and commonsense reasoning.

A second challenge is that commonsense benchmarks are inherently harder to build than knowledge-based test sets. There are no large pools of ready-made questions, so construction typically requires creating scenarios, designing plausible distractors, and obtaining nuanced human judgments. This is further complicated by the fact that commonsense questions often lack a single definitive answer. Meanwhile, widely used benchmarks such as HellaSwag (Zellers et al., 2019), StoryCloze (Mostafazadeh et al., 2016), and CommonsenseQA (Talmor et al., 2019) are saturated, with strong LLMs already scoring above 90 percent. Scaling them with more human annotation is expensive and offers limited headroom for increasing difficulty. Figure 1 illustrates a saturated HellaSwag example. In multilingual settings, the problem is amplified because commonsense is culturally grounded, so questions must be validated to remain fair and meaningful across languages. Existing multilingual commonsense benchmarks (Li et al., 2025;

¹<https://www.anthropic.com/news/claude-4>

Table 1: Comparison of commonsense benchmarks. **NLI** refers to task types where the subsequent content is inferred from the preceding text; these tasks generally preserve the fluency of natural language corpora. **QA** refers to simple question-and-answer formats. **Constructed** refers to more complex, human-defined setups, typically with an instruction. Difficulty is determined by GPT-4 accuracy: * $\geq 80\%$, ** $\geq 50\%$, *** $< 50\%$.

Benchmark	Supported Languages	Total Samples	Local Commonsense	Task Format	Difficulty
HellaSwag	En	10k	✗	NLI	*
StoryCloze	En	1.8k	✗	NLI	*
CommonsenseQA	En	1.1k	✗	QA	*
mCSQA	8	11k	✗	QA	*
CRoW	5	16k	✗	Constructed	*
Com ²	En	3.7k	✗	Constructed	**
HellaSwag-Pro	2	12k (Zh)	✓	Constructed	*
HellaSwagUltra (ours)	61 (5 gold + 56 silver)	163k+	✓	NLI	***

Sakai et al., 2024a; Ismayilzada et al., 2023) cover limited languages and often underrepresent local culture and social context.

A third challenge is difficulty design, which is especially problematic in multilingual settings. Prior work often increases difficulty by changing task formats (Li et al., 2025; Xiong et al., 2025; Ismayilzada et al., 2023), for example causal or multi hop reasoning, abduction, or ordering. In practice, such formats partly measure instruction following rather than core language understanding. Consequently, base or smaller models often collapse under complex formats, while strong models approach saturation, producing a distorted difficulty curve. This reduces the benchmark’s value for guiding pre-training or fine-tuning, since progress appears as abrupt jumps instead of gradual, informative improvements.

To address these challenges, we introduce HellaSwagUltra, a 61-language commonsense reasoning benchmark with a human-verified Gold subset for five languages and LLM-validated silver data for the rest. It is grounded in each language’s native culture, social context, and everyday knowledge. We adopt the natural language inference continuation format used in HellaSwag, because it closely matches the training objective of causal LLMs, supports stable evaluation for base and smaller models, and keeps the instruction for instruction-tuned models minimal. Our automated pipeline first collects Wikipedia and Wikidata pages related to local knowledge in each language, then prompts LLMs to summarize relevant commonsense and generate subtle violations of that knowledge. These commonsense facts are used to define story frameworks and write narrative contexts, with two commonsense intents embedded in each story. In the an-

swer candidates, we inject details that either uphold or subtly contradict the embedded commonsense, forcing models to reason about which continuation is most plausible. This yields fine grained plausibility judgments while avoiding both strong-model saturation and small-model collapse. Table 1 compares HellaSwagUltra with prior work, and Appendix A lists all supported languages.

To summarize, our contributions are as follows:

- We introduce HellaSwagUltra, a large-scale multilingual commonsense reasoning benchmark with over 60 languages and 163k instances, grounded in local cultural knowledge.
- We develop a fully automated construction pipeline that scales across languages while generating realistic scenarios, subtle distractors, and consistent annotations.
- We propose a difficulty scheme that implicitly embeds multiple commonsense facts in each context, enabling stable evaluation for both small and strong models and reducing saturation on advanced models.
- We release HellaSwagUltra-Gold, a human-verified subset for high-stakes evaluation, and report extensive experiments showing that state-of-the-art LLMs still lag far behind human performance.

2 Related Work

Multilingual Benchmarks Existing multilingual benchmarks are broadly of two types. One line translates and extends English benchmarks. BenchMax (Huang et al., 2025) ports a diverse suite into 17 languages but has limited commonsense coverage. MuBench (Han et al., 2025) extends widely

used English pretraining benchmarks, including commonsense and NLU sets such as HellaSwag (Zellers et al., 2019), StoryCloze (Mostafazadeh et al., 2016), SNLI (Bowman et al., 2015), and MultiNLI (Williams et al., 2018), to 61 languages with flexible evaluation formats. However, translated commonsense benchmarks are often saturated, vulnerable to contamination, and may inherit cultural bias. A second line builds native language corpora. INCLUDE (Romanou et al., 2024) collects exam questions in 44 languages to assess local knowledge, and MultiLoKo (Hupkes and Bogoychev, 2025) extracts fact-based QA from multilingual Wikipedia. Both mainly target factual knowledge rather than commonsense reasoning.

Commonsense Reasoning Benchmarks Beyond classic benchmarks such as HellaSwag (Zellers et al., 2019), StoryCloze (Mostafazadeh et al., 2016), and CommonsenseQA (Talmor et al., 2019), recent work increases difficulty via more complex task formats. Com² constructs causal graphs and multiple task types, and HellaSwag Pro (Li et al., 2025) decomposes causal relations into tasks such as backward reasoning and ordering. These designs can reduce textual coherence and yield unstable evaluation, and they often emphasize instruction following over core commonsense understanding. Multilingual commonsense efforts include HellaSwag Pro for Chinese via self bootstrapping and mCSQA (Sakai et al., 2024b) based on ConceptNet alignment, but they cover few languages and provide limited in-depth local commonsense knowledge.

3 HellaSwagUltra

HellaSwagUltra uses the simplest task format: selecting the most plausible continuation given a context. This preserves the semantic coherence of the text and aligns with causal LLM training, making evaluation stable for natural language understanding, including for base and smaller models. The construction pipeline is designed to preserve this simplicity while increasing difficulty through hidden commonsense constraints rather than surface-level task instructions. Construction has three stages: **Knowledge Extraction**, **Structured Commonsense Generation**, and **Test Rollout**. Figure 2 shows the pipeline.

3.1 Knowledge Extraction

To obtain local commonsense knowledge for each language, we filter Wikidata entities² relevant to that language’s local cultural and social background in two steps:

Heuristic Filtering We first filter entities by QID. A predefined type pool excludes broad categories unhelpful for commonsense extraction, such as persons, organizations, geographic locations, and dates; all entities with an instance-of or subclass-of relation to these categories are removed.

LLM-Based Filtering For remaining QIDs, we feed wiki titles and pages to an LLM³, asking whether the article is niche, contains potential commonsense, and carries bias risk. If it passes, we collect its pages in all available languages and prompt the LLM to identify pages tied to each language’s local background.

For each language, we select approximately 1,000 Wiki entities and articles to guide and control the LLM in generating targeted commonsense knowledge.

Wikipedia and Wikidata are only used to identify salient local concepts and contexts. Final questions do not ask models to retrieve encyclopedia facts or name source-page entities; selected concepts become short narratives where models judge whether a continuation violates an unstated commonsense constraint. Thus the target is culturally grounded plausibility reasoning, not factual recall.

3.2 Structured Commonsense Generation

This core stage extracts culturally grounded commonsense for later stories. Using the Wiki pages from the previous step, carefully designed prompts with multiple iterations and validity checks generate structured commonsense with a **Description**, **Premise**, **Consequence**, and **Conflict**.

Description The LLM produces a concise one-sentence description for each commonsense instance, supporting human review and automatic deduplication. We embed descriptions, compute inner products against previous instances, and discard items with excessive similarity.

²https://www.wikidata.org/wiki/Wikidata:Main_Page

³We used Gemini-2.5-Pro throughout the entire data construction process and also involve Claude-Opus-4 for quality control.

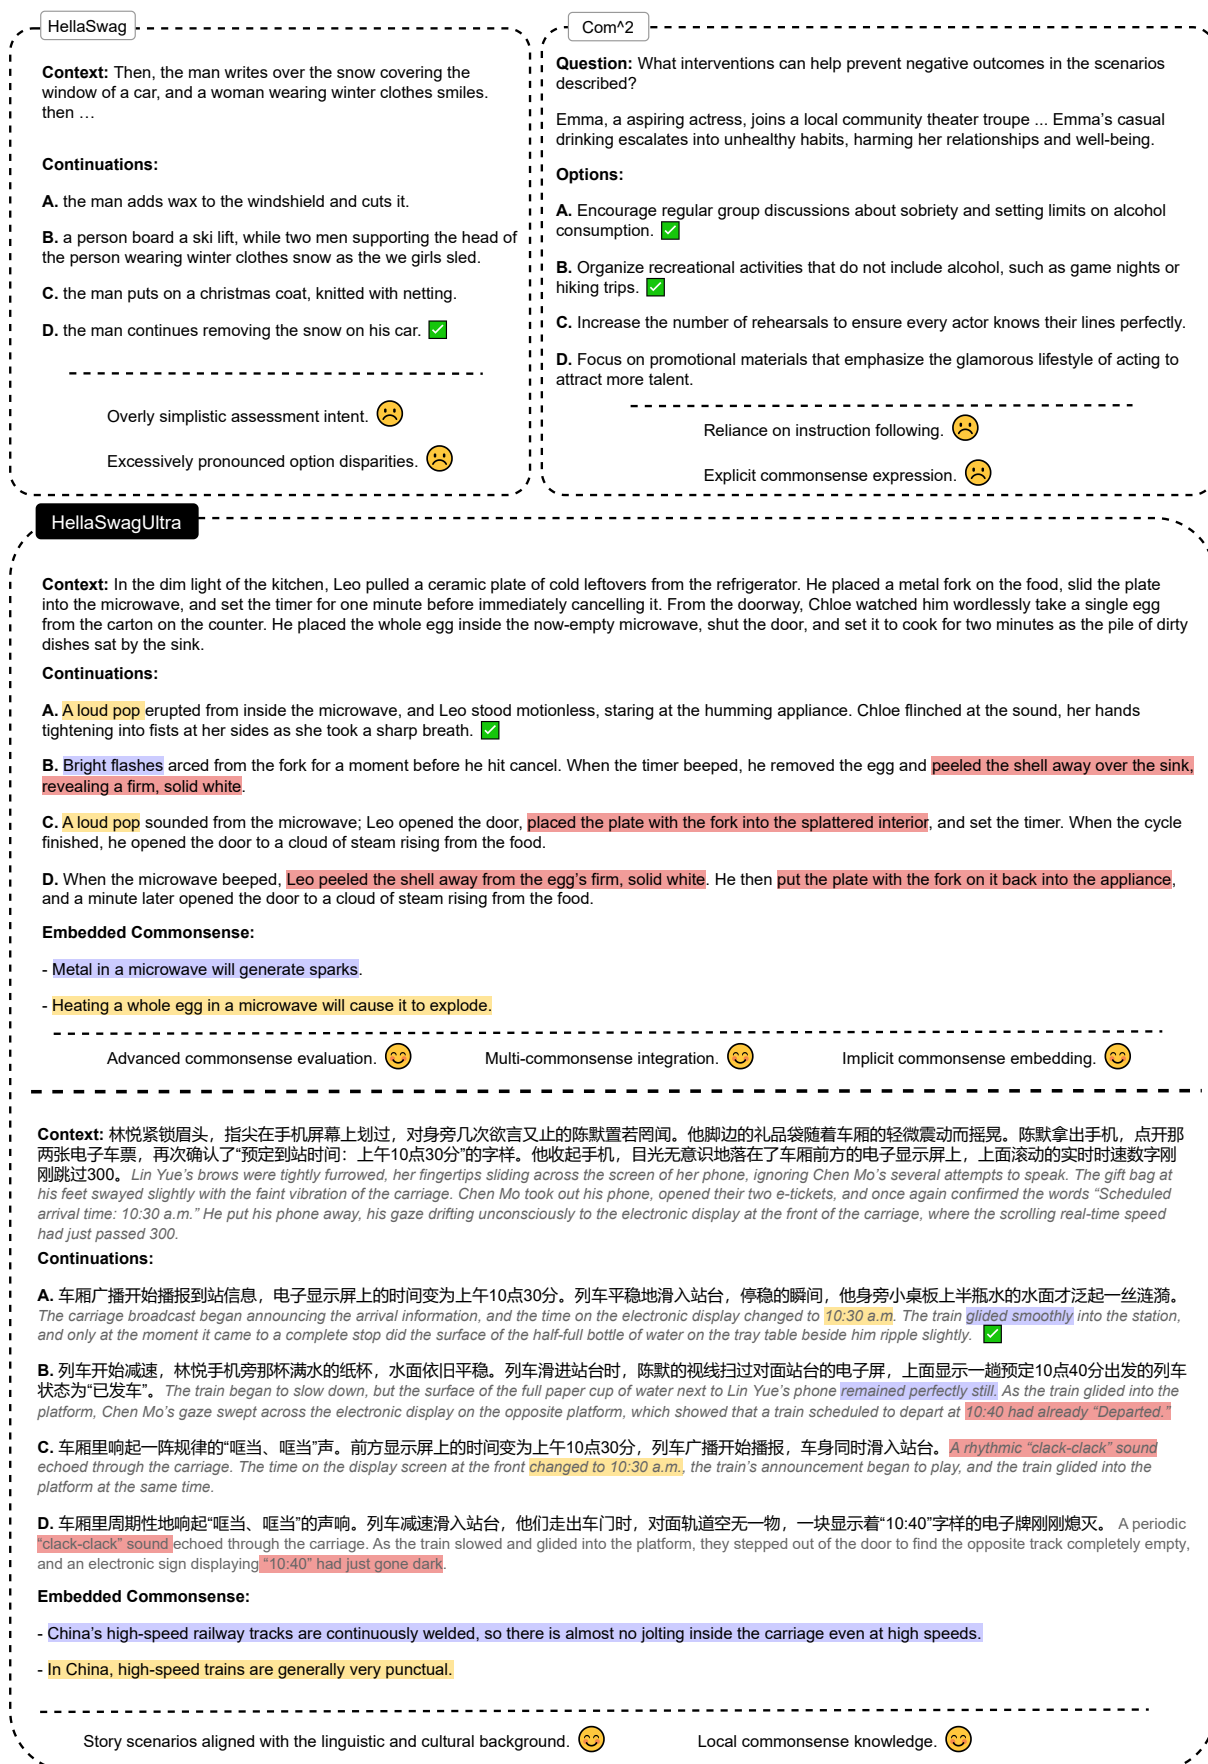


Figure 1: Existing commonsense benchmarks are reaching saturation and cover only a limited set of languages, with insufficient focus on language-specific local commonsense. HellaSwagUltra spans 60+ languages, incorporates a wide range of local, culturally grounded commonsense scenarios, embeds commonsense knowledge implicitly in the context, and offers highly challenging distractor options.

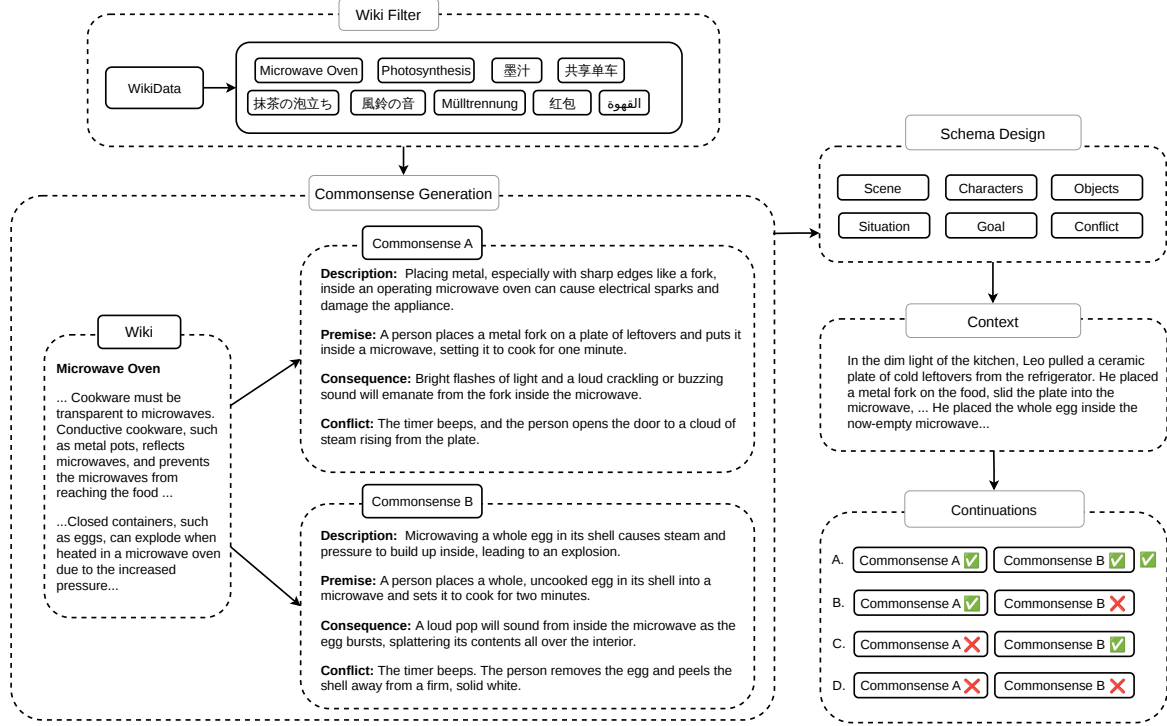


Figure 2: The automatic data construction process consists of three main stages: Knowledge Extraction, Structured Commonsense Generation, and Test Rollout. The Test Rollout stage itself is composed of Schema Design and Context and Continuation Generation.

Premise The LLM creates a premise establishing when the commonsense applies. Generation follows two principles: **Sufficiency and Necessity**: the premise should deterministically lead to a consequence, and the consequence should imply that the premise occurred. **No Outcome Leakage**: the premise must not contain or reveal the consequence.

Consequence Based on the commonsense description and premise, the LLM generates a correct consequence.

Conflict To increase difficulty, the LLM generates a detail that subtly implies a commonsense violation. The principles are: **No Direct Negation**: the conflict must not directly negate people, objects, or events in the premise or consequence. **Objective Description**: it should objectively describe the scene or event, avoid ambiguity or speculation, and exclude subjective thoughts or feelings. **Intrinsic Plausibility**: it must be reasonable and not surreal or impossible.

As Figure 2 shows, each question requires two distinct commonsense instances. After Commonsense A is produced, it becomes a demonstration reference for generating Commonsense B. Beyond

the above principles, B must satisfy **Logical Independence**: whether B is violated should not affect whether A is violated. This avoids overly similar commonsense statements that would weaken challenge and discriminative power. The two-instance design also makes distractors less detectable by shallow lexical cues: an incorrect continuation can remain fluent and locally coherent while quietly violating one embedded constraint.

For quality control, an LLM-based validator checks each requirement using a ten-item checklist compiled from the principles above. For each generated instance, it answers “Yes” if satisfied and “No” with an explanation otherwise. The explanations are logged and incorporated into the next generation prompt, instructing the LLM to revise its output. This process repeats until the instance passes all checklist items.

3.3 Test Rollout

Story Schema Design We do not use commonsense pairs verbatim to generate contexts, since this makes cues obvious, reduces naturalness, and lowers difficulty. Instead, we feed only the two premises and ask the LLM to produce a schema guided by three principles. **Subtle Integration**:

premise details should be woven into the main storyline with minimal exposition, not become the focus. **Completeness**: all premise details must be preserved. **No Outcome Leakage**: the schema must not reveal or hint at consequences or outcomes.

Context and Continuation Generation We then provide the schema and premises to generate the final context under the same principles. Afterward, we combine the consequences and conflicts of Commonsense A and B in different ways and prompt the LLM for continuations. The option with both consequences is correct; any option containing a conflict from either commonsense is a distractor.

Quality Control and Annotation In addition to commonsense validation, we quality-check final items. To reduce random guessing, we resample each distractor twice, yielding ten candidate options per question. We then give multiple LLMs the full question plus explicit commonsense hints; if they fail to select the correct answer, we discard the item. This filter is intentionally conservative: if models with the intended commonsense hints cannot recover the correct option, the item is likely ambiguous, underspecified, or noisy rather than genuinely difficult. We also ask an LLM to annotate local relevance. **Local Background** marks items whose context contains target-language-specific cultural, social, or environmental elements, such as local objects, events, institutions, or routines, though they may remain solvable with general commonsense. **Local Commonsense** is stricter: solving requires culture-specific knowledge from the language region, not just recognizing a local surface detail. The labels can overlap, but local background alone does not imply local commonsense. Prompts are in Appendix B.

3.4 Human Evaluation

We recruited language-specific annotators, with at least three per language. For human performance evaluation, we sampled 100 questions for each of English, Chinese, Arabic, German, French, Portuguese, and Indonesian. Annotators averaged 92% accuracy on 50 questions with underlying commonsense provided, and 73% on 50 questions without hints. Beyond this evaluation, we fully verify and correct items for correctness. HellaSwagUltra-Gold contains 2,926 verified items across English, German, Arabic, French, and Indonesian (Table 2).

Each Gold item is checked by two annotators; disagreements go to a third annotator for adjudication and final consensus. Given the scale of HellaSwagUltra, we will maintain and update the dataset long term. Annotation and cost details are in Appendix C.

3.5 Statistics

Table 2: Statistics of HellaSwagUltra and HellaSwagUltra-Gold. LB: Local Background; LC: Local Commonsense. LB and LC counts use each language’s full set, and the two labels may overlap.

Language	Total	Verified	LB	LC
ALL	163,138	2,926	112,784	77,713
<i>Gold languages</i>				
EN	2,943	766	699	558
DE	2,863	605	1,095	1,128
AR	2,820	652	1,561	1,235
FR	2,816	568	1,115	1,119
ID	2,946	335	2,274	1,417

Table 2 reports HellaSwagUltra statistics. Of all 163,138 samples, 112,784 (69.1%) have Local Background labels and 77,713 (47.6%) have Local Commonsense labels. These percentages use the full-set denominator; LB and LC can overlap. The language rows report full-set LB and LC counts alongside the number of Gold items. Among the five Gold languages, English has the lowest full-set proportion of both labels, consistent with its widespread use across diverse regions, which makes its content more likely to be perceived as general rather than culturally anchored. German and French show similarly low local-commonsense proportions, reflecting closer cultural affinity with English and shared background knowledge. Arabic has a much higher proportion of local-commonsense items, highlighting the distinct cultural specificity captured in this subset.

4 Evaluation

4.1 Setup

Task Format HellaSwagUltra follows HellaSwag’s core format: choose the most plausible continuation given a context. This preserves natural text, aligns with causal LLM training, and enables stable evaluation. For base models, we use Cloze-style scoring (Clark et al., 2018): compute each candidate’s perplexity and select the lowest. For instruction-tuned models, we use the multiple-choice prompt “Which option is the most plausible continuation?” in both English and localized form,

Table 3: Model performance on HellaSwagUltra. **ALL**: average accuracy across languages. **HIGH/MID/LOW**: high-, mid-, and low-resource averages. **LB**: local-background items; **LC**: culture-specific commonsense items; **GC**: language-agnostic commonsense. **VERIFIED**: HellaSwagUltra-Gold. Base models use Cloze scoring, so base and instruction-tuned rows should be compared mainly within protocol.

Model	ALL	HIGH	MID	LOW	LB	LC	GC	VERIFIED
<i>Base Models</i>								
Qwen3-14B-Base	43.86	48.45	45.41	40.47	42.94	43.08	44.15	47.68
Qwen2.5-14B	41.12	48.17	42.69	36.69	40.65	40.17	41.61	47.69
Qwen2.5-32B	42.25	50.53	43.58	37.52	41.58	41.52	42.55	49.34
Qwen2.5-72B	44.84	52.53	47.34	39.28	43.92	43.71	45.22	48.85
gemma-3-12b-pt	47.77	48.09	49.92	45.67	47.08	46.61	48.27	47.51
gemma-3-27b-pt	50.88	51.06	53.12	48.73	50.42	50.09	51.28	49.26
gemma-2-9b	43.94	46.44	45.98	41.00	43.00	43.04	44.16	47.20
gemma-2-27b	47.15	49.25	50.24	43.41	46.13	45.76	47.63	48.56
<i>Instruct Models</i>								
Qwen2.5-14B-Instruct	37.54	46.13	38.94	32.61	38.21	38.52	37.25	50.22
Qwen2.5-32B-Instruct	39.59	47.61	41.41	34.52	40.29	41.13	39.19	50.22
Qwen2.5-72B-Instruct	40.18	47.76	41.42	35.82	40.29	40.45	40.00	49.47
gemma-3-12b-it	34.57	36.27	35.18	33.29	35.34	34.68	34.30	40.84
gemma-3-27b-it	42.02	43.55	41.66	41.70	42.19	42.28	41.94	48.32
gemma-2-9b-it	31.51	33.45	31.21	30.98	32.27	32.77	30.89	37.31
gemma-2-27b-it	37.26	40.28	37.75	35.54	37.50	37.64	36.98	45.01
<i>Proprietary Model</i>								
GPT-4o	40.34	47.17	41.87	36.02	40.29	40.11	40.40	50.64
Claude Sonnet 4	63.12	65.58	64.46	60.86	62.78	62.11	63.52	60.14
Claude Opus 4	64.53	66.28	65.31	63.07	63.75	63.06	65.11	59.57
Gemini-2.5-Pro	72.13	70.31	73.11	72.00	71.71	71.30	72.60	62.53




Figure 3: Model performance with and without thinking. Blue bars denote direct evaluation without thinking or chain-of-thought prompting; orange bars denote thinking mode for Claude models and chain-of-thought prompting for GPT-4o.

and report localized-instruction results to reflect multilingual performance. Because the protocols target different interfaces, we interpret base and instruction-tuned results mainly within each group, not as absolute cross-paradigm rankings.

All evaluations are deterministic: base models use candidate perplexity scoring without sampling, while instruction-tuned models decode with temperature 0.

Metric We report accuracy for base models (lowest-PPL choice) and Exact Match for instruction-tuned models (generated answer matches the correct option).

Models We evaluate multilingual open-source base models from the Qwen (Qwen et al., 2025;

Yang et al., 2025) and Gemma families (Team et al., 2024b, 2025), together with their instruction-tuned variants. As noted above, base and instruction-tuned protocols differ. We also benchmark capable closed-source models, including GPT-4o (OpenAI et al., 2024), Claude Opus 4, and Claude Sonnet 4. All models are evaluated without task-specific fine-tuning, since the goal is benchmark construction and analysis rather than adaptation recipes.

4.2 Results

Table 3 shows that all evaluated models perform suboptimally on HellaSwagUltra, indicating sufficient challenge and reduced saturation. Importantly, the difficulty comes from stronger commonsense reasoning demands rather than com-

Gemini data -	ALL	HIGH	ALL	HIGH
	72.1	70.3	64.5	66.3
	MID	LOW	MID	LOW
	73.1	72.0	65.3	63.1
Opus data -	ALL	HIGH	ALL	HIGH
	55.2	58.0	47.0	54.6
	MID	LOW	MID	LOW
	57.5	52.0	51.7	39.4
	Gemini eval		Opus eval	

Figure 4: Cross-evaluation of datasets generated by different models. Rows identify the data generator and columns identify the evaluator; each generator contributes 300 samples per language, and cells report accuracy (%).

plex instructions or evaluation metrics. Within each family, larger models consistently perform better. Across languages, all models show a noticeable drop on low-resource languages. Among open-source families, Gemma is relatively balanced, while Qwen shows a larger high- versus low-resource gap. As expected, closed-source models generally outperform open-source ones, and GPT-4o, which is not a reasoning model, lags substantially behind Claude models.

Gemini scores far higher on the full set, but the gap shrinks on the human-verified subset. This suggests generator- or style-specific effects may contribute and underscores the importance of human verification. At the same time, the verified results show that the benchmark remains difficult even after human correction, so the challenge is not merely annotation noise. Silver data for languages not yet human-verified remain valuable because they provide broad diagnostic coverage, while Gold data anchors claims that require stronger correctness guarantees.

4.3 Effect of Thinking

We evaluate thinking mode on HellaSwagUltra. For Opus and Sonnet, we compare thinking enabled and disabled; for GPT-4o, which is not a reasoning model by default, we use a Chain-of-Thought prompt before answer generation and compare it with direct output. Figure 3 shows that all three models gain substantially on the full dataset with thinking, while gains are smaller on HellaSwagUltra-Gold, with Opus slightly de-

clining. On the local-commonsense subset, Sonnet and GPT-4o gain more than on the non-local-commonsense subset.

4.4 Potential Model Bias Study

To test whether the generator LLM gains an advantage on its own data, we used Claude Opus 4 to generate 300 samples per language with the same pipeline, then evaluated Gemini-2.5-Pro and Claude Opus 4 on both generators’ data. Figure 4 shows that Gemini’s ALL-score margin over Opus is 7.6 points on Gemini-generated data ($72.1 - 64.5$) and 8.2 points on Opus-generated data ($55.2 - 47.0$). Thus, Gemini’s relative advantage is not larger on its own data, and Opus does not gain on its self-generated data. This pattern argues against a simple self-generator advantage, but it does not rule out residual generator or style bias. We therefore treat generator effects as a source of noise to monitor: full-set scores support broad language-level diagnosis, while Gold scores should anchor absolute capability claims and comparisons requiring stronger validity.

5 Conclusion

Faced with the saturation of commonsense reasoning benchmarks and the scarcity of multilingual resources, this paper introduces HellaSwagUltra, a new dataset supporting over 60 languages, focused on challenging commonsense reasoning and language local knowledge. The results show that HellaSwagUltra is highly challenging and strongly discriminative across models, filling an important gap in multilingual evaluation of local commonsense for base models. By preserving a simple continuation format, HellaSwagUltra evaluates multilingual plausibility reasoning without relying on complex task instructions, while its automated pipeline supports continued expansion. The human-verified Gold subset further provides a reliable anchor for higher-confidence evaluation and future benchmark maintenance, while the silver split enables broad diagnostic coverage for languages where full expert verification is still costly. Given the wide language coverage, we plan to maintain and continuously update the human-verified branch, HellaSwagUltra-Gold, until it achieves broader coverage.

Limitations

HellaSwagUltra is built with an automated pipeline and relies on a single strong LLM for generation

and several LLM-based checks. This can introduce model-specific style artifacts and latent cultural biases, especially for low-resource languages where Wikipedia coverage is sparse and local knowledge is under-documented. While we mitigate these risks with validation and targeted human review, the current human-verified subset covers only a small fraction of the full benchmark, so some annotation noise and ambiguous items may remain, particularly in silver data for languages not yet covered by Gold verification. In addition, our NLI continuation format focuses on plausibility judgments and may not capture other forms of commonsense reasoning, such as interactive planning or grounded decision making.

Future work will expand HellaSwagUltra-Gold to more languages and larger coverage, add deeper expert verification for local commonsense labels, and diversify knowledge sources beyond Wikipedia to better represent everyday culture. We also plan to develop stronger bias and contamination audits, including checks for generator-specific artifacts and source coverage imbalance, and to explore complementary evaluation formats that preserve stable difficulty while testing a broader range of reasoning behaviors. These additions will help separate failures caused by missing cultural grounding from those caused by language modeling, generation artifacts, or format sensitivity, and will make future releases easier to audit across Gold data, silver data, and pipeline revisions. This separation is especially important for low-resource languages, where sparse sources and automatic generation artifacts can otherwise be difficult to disentangle.

References

- Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. 2015. [A large annotated corpus for learning natural language inference](#). In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 632–642, Lisbon, Portugal. Association for Computational Linguistics.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. [Think you have Solved Question Answering? Try ARC, the AI2 Reasoning Challenge](#). *Preprint*, arXiv:1803.05457.
- Wenhan Han, Yifan Zhang, Zhixun Chen, Binbin Liu, Haobin Lin, Bingni Zhang, Taifeng Wang, Mykola Pechenizkiy, Meng Fang, and Yin Zheng. 2025. [MuBench: Assessment of Multilingual Capabilities of Large Language Models Across 61 Languages](#). *Preprint*, arXiv:2506.19468.
- Xu Huang, Wenhao Zhu, Hanxu Hu, Conghui He, Lei Li, Shujian Huang, and Fei Yuan. 2025. [BenchMAX: A Comprehensive Multilingual Evaluation Suite for Large Language Models](#). *Preprint*, arXiv:2502.07346.
- Dieuwke Hupkes and Nikolay Bogoychev. 2025. [MultiLoKo: A multilingual local knowledge benchmark for LLMs spanning 31 languages](#). *Preprint*, arXiv:2504.10356.
- Mete Ismayilzada, Debjit Paul, Syrielle Montariol, Mor Geva, and Antoine Bosselut. 2023. [CROW: Benchmarking Commonsense Reasoning in Real-World Tasks](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9785–9821, Singapore. Association for Computational Linguistics.
- Viet Lai, Chien Nguyen, Nghia Ngo, Thuat Nguyen, Franck Dernoncourt, Ryan Rossi, and Thien Nguyen. 2023. [Okapi: Instruction-tuned Large Language Models in Multiple Languages with Reinforcement Learning from Human Feedback](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 318–327, Singapore. Association for Computational Linguistics.
- Haonan Li, Yixuan Zhang, Fajri Koto, Yifei Yang, Hai Zhao, Yeyun Gong, Nan Duan, and Timothy Baldwin. 2024. [CMMLU: Measuring massive multi-task language understanding in Chinese](#). *Preprint*, arXiv:2306.09212.
- Xiaoyuan Li, Moxin Li, Rui Men, Yichang Zhang, Keqin Bao, Wenjie Wang, Fuli Feng, Dayiheng Liu, and Junyang Lin. 2025. [HellaSwag-Pro: A Large-Scale Bilingual Benchmark for Evaluating the Robustness of LLMs in Commonsense Reasoning](#). *Preprint*, arXiv:2502.11393.
- Nasrin Mostafazadeh, Nathanael Chambers, Xiaodong He, Devi Parikh, Dhruv Batra, Lucy Vanderwende, Pushmeet Kohli, and James Allen. 2016. [A Corpus and Evaluation Framework for Deeper Understanding of Commonsense Stories](#). *Preprint*, arXiv:1604.01696.
- OpenAI, Aaron Hurst, Adam Lerer, Adam P. Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, A. J. Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, Aleksander Madry, Alex Baker-Whitcomb, Alex Beutel, Alex Borzunov, Alex Carney, Alex Chow, Alex Kirillov, Alex Nichol, and 400 others. 2024. [GPT-4o System Card](#). *Preprint*, arXiv:2410.21276.
- Qwen, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, and 24 others. 2025. [Qwen2.5 Technical Report](#). *Preprint*, arXiv:2412.15115.

- Angelika Romanou, Negar Foroutan, Anna Sotnikova, Zeming Chen, Sree Harsha Nelaturu, Shivalika Singh, Rishabh Maheshwary, Micol Altomare, Mohamed A. Haggag, Snegha A, Alfonso Amayuelas, Azril Hafizi Amirudin, Viraat Aryabumi, Danylo Boiko, Michael Chang, Jenny Chim, Gal Cohen, Aditya Kumar Dalmia, Abraham Diress, and 40 others. 2024. [INCLUDE: Evaluating Multilingual Language Understanding with Regional Knowledge](#). *Preprint*, arXiv:2411.19799.
- Yusuke Sakai, Hidetaka Kamigaito, and Taro Watanabe. 2024a. [mCSQA: Multilingual Commonsense Reasoning Dataset with Unified Creation Strategy by Language Models and Humans](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 14182–14214, Bangkok, Thailand. Association for Computational Linguistics.
- Yusuke Sakai, Hidetaka Kamigaito, and Taro Watanabe. 2024b. [mCSQA: Multilingual Commonsense Reasoning Dataset with Unified Creation Strategy by Language Models and Humans](#). *Preprint*, arXiv:2406.04215.
- Shivalika Singh, Angelika Romanou, Cl  mentine Fourrier, David I. Adelani, Jian Gang Ngui, Daniel Vila-Suero, Peerat Limkonchotiwat, Kelly Marchisio, Wei Qi Leong, Yosephine Susanto, Raymond Ng, Shayne Longpre, Wei-Yin Ko, Sebastian Ruder, Madeline Smith, Antoine Bosselut, Alice Oh, Andre F. T. Martins, Leshem Choshen, and 5 others. 2025. [Global MMLU: Understanding and Addressing Cultural and Linguistic Biases in Multilingual Evaluation](#). *Preprint*, arXiv:2412.03304.
- Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. [CommonsenseQA: A Question Answering Challenge Targeting Commonsense Knowledge](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4149–4158, Minneapolis, Minnesota. Association for Computational Linguistics.
- Gemini Team, Rohan Anil, Sebastian Borgeaud, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M. Dai, Anja Hauth, Katie Millican, David Silver, Melvin Johnson, Ioannis Antonoglou, Julian Schrittwieser, Amelia Glaese, Jilin Chen, Emily Pitler, Timothy Lillicrap, Angeliki Lazaridou, and 1331 others. 2024a. [Gemini: A Family of Highly Capable Multimodal Models](#). *Preprint*, arXiv:2312.11805.
- Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ram  , Morgane Riv  re, Louis Rouillard, Thomas Mesnard, Geoffrey Cideron, Jean-bastien Grill, Sabela Ramos, Edouard Yvinec, Michelle Casbon, Etienne Pot, Ivo Penchev, and 197 others. 2025. [Gemma 3 Technical Report](#). *Preprint*, arXiv:2503.19786.
- Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, L  onard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ram  , Johan Ferret, Peter Liu, Pouya Tafti, Abe Friesen, Michelle Casbon, Sabela Ramos, Ravin Kumar, Charline Le Lan, Sammy Jerome, and 177 others. 2024b. [Gemma 2: Improving Open Language Models at a Practical Size](#). *Preprint*, arXiv:2408.00118.
- Adina Williams, Nikita Nangia, and Samuel Bowman. 2018. [A Broad-Coverage Challenge Corpus for Sentence Understanding through Inference](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1112–1122, New Orleans, Louisiana. Association for Computational Linguistics.
- Kai Xiong, Xiao Ding, Yixin Cao, Yuxiong Yan, Li Du, Yufei Zhang, Jinglong Gao, Jiaqian Liu, Bing Qin, and Ting Liu. 2025. [Com²: A Causal-Guided Benchmark for Exploring Complex Commonsense Reasoning in Large Language Models](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16119–16140, Vienna, Austria. Association for Computational Linguistics.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, and 41 others. 2025. [Qwen3 Technical Report](#). *Preprint*, arXiv:2505.09388.
- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. 2019. [HellaSwag: Can a Machine Really Finish Your Sentence?](#) In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4791–4800, Florence, Italy. Association for Computational Linguistics.

A Language Coverage

Table 4 presents the language covered by HellaSwagUltra. Considering only native speakers, these languages cover over 60% of the global population. When including second-language speakers, the coverage exceeds 99% worldwide.

B Prompts

The prompts used in data collection are as follows.

C Human Evaluation and Cost

We recruited human annotators who were required to hold at least a college degree, demonstrate C1-

Table 4: Languages sorted by native speakers and ratios in Common Crawl (HIGH at left, MID center, LOW right)

Code	Name	Speakers	Tokens	Code	Name	Speakers	Tokens	Code	Name	Speakers	Tokens
zh	Chinese	1390M	6.34%	vi	Vietnamese	86M	1.35%	hi	Hindi	345M	0.31%
es	Spanish	484M	4.14%	tr	Turkish	85M	0.98%	bn	Bengali	242M	0.18%
ar	Arabic	411M	0.78%	ms	Malay	82M	0.03%	mr	Marathi	83M	0.04%
en	English	390M	42.62%	ur	Urdu	78M	0.04%	te	Telugu	83M	0.03%
pt	Portuguese	250M	1.51%	id	Indonesian	75M	1.05%	ta	Tamil	79M	0.09%
ru	Russian	145M	9.16%	fa	Persian	65M	0.79%	jv	Javanese	69M	0.00%
ja	Japanese	124M	4.72%	pl	Polish	38M	1.69%	gu	Gujarati	58M	0.03%
ko	Korean	81M	0.84%	th	Thai	38M	0.64%	my	Burmese	33M	0.03%
de	German	76M	5.21%	uk	Ukrainian	32M	0.60%	pa	Punjabi	32M	0.01%
fr	French	74M	4.10%	ro	Romanian	24M	0.64%	tl	Tagalog	28M	0.02%
it	Italian	63M	2.33%	nl	Dutch	23M	1.57%	uz	Uzbek	27M	0.01%
				el	Greek	12M	0.69%	az	Azerbaijani	24M	0.10%
				bg	Bulgarian	8M	0.32%	ceb	Cebuano	21M	0.00%
				hr	Croatian	5.1M	0.24%	sw	Swahili	16M	0.01%
				sk	Slovak	5M	0.35%	km	Khmer	16M	0.02%
				he	Hebrew	5M	0.27%	sq	Albanian	7.5M	0.05%
				lt	Lithuanian	2.8M	0.18%	af	Afrikaans	7M	0.01%
				lv	Latvian	1.75M	0.10%	no	Norwegian	5.3M	0.37%
				et	Estonian	1.1M	0.14%	da	Danish	5M	0.36%
				sv	Swedish	10M	0.63%	fi	Finnish	5M	0.41%
				cs	Czech	11M	1.02%	is	Icelandic	0.314M	0.04%
				hu	Hungarian	13M	0.49%	ga	Irish	—	0.01%
				sr	Serbian	9M	0.21%	ca	Catalan	4M	0.17%
				sl	Slovenian	2.1M	0.13%	kk	Kazakh	15M	0.04%
								kn	Kannada	44M	0.01%
								ml	Malayalam	38M	0.02%

level English proficiency (or an equivalent certification), and be native speakers of the languages they were assigned to evaluate. Annotators were paid at an hourly rate of \$16, with a maximum of 8 working hours per day. To date, the total cost of human annotation is approximately \$19,200.

The human performance evaluation and Gold verification are separate procedures. The performance study samples 100 questions per language across seven languages, while HellaSwagUltra-Gold contains fully corrected items for English, German, Arabic, French, and Indonesian. In Gold verification, each annotator provides a complete three-field label set: the correct option, Local Background, and Local Commonsense. The first two annotators have 76% pooled raw agreement across the five Gold languages under an exact-match criterion that requires all three fields to match. When their complete label sets differ, a third annotator adjudicates the item. If the third annotation matches either of the first two complete label sets, that two-of-three label set is retained; if it matches neither, the item is discarded. Among items sent to adjudication, 21% received three different label sets and were discarded.

In addition, the API cost for LLM calls during data collection is approximately \$36,800.

Commonsense Generation

We are designing a local commonsense knowledge in {LANGUAGE} that will be subtly implied in a story. The `conflict_detail` is a subtle narrative detail that may hint that the commonsense has been violated. Later, this will be used to generate an incorrect story continuation.

[Reference Wiki]

{WIKITEXT}

[TASK]

Invent a new piece of local commonsense knowledge in {LANGUAGE}.

Return a single JSON object using this exact schema with the content in {LANGUAGE}:

concept : short noun phrase (1–3 words)

description : one sentence explaining the concept

premise : a brief and typical sign or condition under which the consequence usually holds (no need for full coverage)

- The premise is a specific condition or context.
- The premise should **not include or imply the consequence itself** — it must be a distinct and self-contained condition, not a restatement or soft prediction of the result.
- The premise must be both necessary **and** sufficient for the consequence.

consequence : what is expected under the premise

conflict_detail : a subtle, objective event or scene implying the consequence may have been violated (do NOT state the violation or directly revise the property)

- `conflict_detail` should hint at a breach indirectly and be subtle.
- Describe only observable facts, actions, or physical details.
- Do **NOT** mention the items or properties that appear in the consequence.
- Do **NOT** overemphasize or elaborate on the conflicting detail.
- Do **NOT** emphasize or mention what did *not* happen.
- Do **NOT** justify, rationalize, or explain the detail that contradicts commonsense.
- Do **NOT** use negation words (e.g., “not”, “no”, “never”, “didn’t”, “hasn’t”, “without”, “failed to”).
- Do **NOT** use speculative or ambiguous expressions (e.g., “seems”, “appears”, “perhaps”, “maybe”, “apparently”, “as if”).
- Do **NOT** use contrastive words like “but”, “however”.
- Do **NOT** include thoughts, feelings.

Quality requirements:

The commonsense must be strong and widely accepted: in ordinary contexts it should hold with certainty; violating it should make the scenario feel blatantly unrealistic or jarringly wrong to a knowledgeable reader.

Given the premise, the `conflict_details` should be virtually impossible to happen in reality.

Ensure **diversity** from the reference examples, the topic **MUST BE DIFFERENT** from any of the reference examples.

Consider diverse types and maintain a numerical balance between **traditional and modern, cultural and scientific, local and worldwide** commonsenses.

Both the premise and the consequence should have a **moderate level of abstraction**: avoid overly specific names, locations, or one-time events. The statements should apply to many plausible real-world situations.

Reference examples:

{examples}

Commonsense Revision

We are designing a local commonsense knowledge in {LANGUAGE} that will be subtly implied in a story. The `conflict_detail` is a subtle narrative detail that may hint that the commonsense has been violated. Later, this will be used to generate an incorrect story continuation.

[REFERENCE EXAMPLES]

{examples}

[PREVIOUS VERSION]

{previous_json_pretty}

[VALIDATOR COMMENTS]

{feedback_rules_block}

[TASK]

Please revise the above commonsense item to better follow the rules and address the validator feedback.

Return a single JSON object using this exact schema with the content in {LANGUAGE}:

concept : short noun phrase (1–3 words)

description : one sentence explaining the concept

premise : a brief and typical sign or condition under which the consequence usually holds (no need for full coverage)

- The premise is a specific condition or context.
- The premise should **not include or imply the consequence itself** — it must be a distinct and self-contained condition, not a restatement or soft prediction of the result.
- The premise must be both necessary **and** sufficient for the consequence.

consequence : what is expected under the premise

conflict_detail : a subtle, objective event or scene implying the consequence may have been violated (do NOT state the violation or directly revise the property)

- `conflict_detail` should hint at a breach indirectly and be subtle.
- Describe only observable facts, actions, or physical details.
- Do **NOT** mention the items or properties that appear in the consequence.
- Do **NOT** overemphasize or elaborate on the conflicting detail.
- Do **NOT** emphasize or mention what did not happen.
- Do **NOT** justify, rationalize, or explain the detail that contradicts commonsense.
- Do **NOT** use negation words (e.g., 'not', 'no', 'never', 'didn't', 'hasn't', 'without', 'failed to').
- Do **NOT** use speculative or ambiguous expressions (e.g., 'seems', 'appears', 'perhaps', 'maybe', 'apparently', 'as if').
- Do **NOT** use contrastive words like 'but', 'however'.
- Do **NOT** include thoughts, feelings.

Quality requirements:

The commonsense must be strong and widely accepted: in ordinary contexts it should hold with certainty; violating it should make the scenario feel blatantly unrealistic or jarringly wrong to a knowledgeable reader.

Given the premise, the `conflict_details` should be virtually impossible to happen in reality.

Ensure **diversity** from the reference examples, the topic **MUST BE DIFFERENT** from any of the reference examples.

Consider diverse types and maintain a numerical balance between **traditional and modern, cultural and scientific, local and worldwide** commonsenses.

Both the premise and the consequence should have a **moderate level of abstraction**: avoid overly specific names, locations, or one-time events. The statements should apply to many plausible real-world situations.

Commonsense Validator

We are designing a **{LANGUAGE} commonsense knowledge** item that will be subtly implied in a story. 'description' explains the piece of commonsense knowledge. 'premise' and 'consequence' denote, respectively, the preconditions under which this commonsense holds and its expected result. 'conflict_detail' is a subtle detail that hints at a violation of that result. Your job is to evaluate whether the defined commonsense knowledge and the 'conflict_detail' are valid and well-formed for this purpose based on the following rule.

[RULE]

{RULE_TEXT}

[GUIDELINE]

{GUIDELINE_TEXT}

[COMMONSENSE UNDER REVIEW]

{COMMONSENSE}

Respond in JSON with:

```
{
  "comment": string,
  "decision": "pass" | "fail"
}
```

Story Schema Builder

We are designing a **beginning** of a short realistic story in {LANGUAGE} that integrates both a visible storyline and a hidden layer.

You are given some subtle details.

Your task is to design the schema.

Important constraints:

- The subtle details must **NOT** be the focus and mainstream of the story, nor the characters' main activity.
- Subtly weave the subtle details into the main storyline with minimal exposition.
- Ensure the information in the provided details is **accurately** and **completely** included.
- Do not reveal or speculate about the continuation of the subtle details.

Please provide the following fields:

1. **Scene:** Where does the story take place? Describe the physical and social setting briefly.
2. **Characters:** List 2–3 people involved in the scene, with their names, roles, motivations, features, characteristics, relationship, etc.
3. **Objects:** Any key items or tools present in the scene.
4. **Situation description:** A short paragraph (3–4 sentences) describing what's happening, *without stating* the commonsense.
5. **Goal or activity:** What is the apparent goal of the characters in the scene?
6. **Visible tension or obstacle:** Is there any small conflict or uncertainty that drives the scene forward?

You can add more fields.

Return all fields in plain text in {LANGUAGE}.

Example

Subtle details: Two people are sitting in a café and talking.

Output in English:

Scene: A quiet café in the afternoon.

Characters: Lisa, a journalist; Mark, her childhood friend.

Objects: Coffee cups, a notepad, a scarf hanging behind Lisa's chair.

Situation description: Lisa leans across the table, her voice low as she asks Mark about the article. He listens, occasionally glancing at the entrance. The café hums softly around them.

Goal or activity: They are catching up and discussing a sensitive interview.

Visible tension or obstacle: Lisa is worried someone may overhear them.

Now you generate:

Subtle details:

{premise}

Output in {LANGUAGE}:

Story Context Generator

Given a designed schema, write a beginning of a short, realistic story in {LANGUAGE}.

Constraints:

- The story **must include natural and appropriate mentions** of all people, objects, or situations referenced in the provided schema, but they should appear **organically and with narrative motivation** — not feel forced or inserted just to match the fact.
- The tone should be grounded, realistic, and coherent.
- The story should reflect the described **scene, characters, objects, activity, and visible tension** through concrete actions, sensory details, or dialogue.
- No explaining or summarizing; let the details emerge naturally.
- Subtly weave the provided subtle details into the main storyline with minimal exposition. Ensure all the key information is **accurately** and **completely** included.
- The embedded subtle details must **NOT** be the focus and mainstream of the story, nor the characters' main activity.
- Focus on objective, observable descriptions of actions, settings, and dialogue.
- Do **NOT** add thoughts, feelings, or commentary.
- Do **NOT** mention what did **not** happen.
- Do **NOT** use speculative or ambiguous expressions (e.g., “seems”, “appears”, “perhaps”, “maybe”, “apparently”, “as if”).
- The story should be in 4–5 sentences.

Example

Schema:

Scene: A quiet café in the afternoon.

Characters: Lisa, a journalist; Mark, her childhood friend.

Objects: Coffee cups, a notepad, a scarf hanging behind Lisa's chair.

Situation description: Lisa leans across the table, her voice low as she asks Mark about the article. He listens, occasionally glancing at the entrance. The café hums softly around them.

Goal or activity: They are catching up and discussing a sensitive interview.

Visible tension or obstacle: Lisa is worried someone may overhear them.

Subtle details:

Two people are sitting in a café and talking.

Story in English:

Lisa lowered her voice, scribbling something in her notepad as Mark leaned in. The scarf behind her chair fluttered slightly as the door opened. He looked up, eyes scanning the new arrival. “Do you think they’re listening?” she whispered.

Now complete the story based on the following schema:

{schema}

Subtle details:

{details}

Story in {LANGUAGE}:

Story Continuation A (Positive)

You are given the schema of a short story in {LANGUAGE}, a beginning and some subtle details.
Your task is to write a plausible continuation of the story:

- The continuation must naturally **follow from the story so far**, not repeat or revise the story.
- Subtly weave the follow-up of the provided details into the main storyline with minimal exposition.
- Do not justify, rationalize, or explain the details.
- Focus on objective, observable descriptions of actions, settings, and dialogue.
- Do **NOT** add thoughts, feelings, or commentary.
- Do **NOT** mention events or outcomes that did **not** happen — focus on what is occurring in the scene.
- Do **NOT** use negation words (e.g., “not”, “no”, “never”, “didn’t”, “hasn’t”, “without”, “failed to”).
- Do **NOT** use speculative or ambiguous expressions (e.g., “seems”, “appears”, “perhaps”, “maybe”, “apparently”, “as if”).
- The continuation should be in **1–2 sentences**.

Example

Schema:

Scene: A quiet café in the afternoon.

Characters: Lisa, a journalist; Mark, her childhood friend.

Objects: Coffee cups, a notepad, a scarf hanging behind Lisa’s chair.

Situation description: Lisa leans across the table, her voice low as she asks Mark about the article. He listens, occasionally glancing at the entrance. The café hums softly around them.

Goal or activity: They are catching up and discussing a sensitive interview.

Visible tension or obstacle: Lisa is worried someone may overhear them.

Story so far:

Lisa lowered her voice, scribbling something in her notepad as Mark leaned in. The scarf behind her chair fluttered slightly as the door opened. He looked up, eyes scanning the new arrival. “Do you think they’re listening?” she whispered.

Subtle details:

One cannot see what is happening behind himself/herself.

Continuation in English:

Mark shook his head subtly, his eyes drifting past Lisa toward the window. A delivery man stepped inside, pausing to check the receipt in his hand.

Now you generate

Schema:

{schema}

Story so far:

{story}

Subtle details:

{details}

Continuation in {LANGUAGE}:

Story Continuation B (Negative)

You are given the schema of a short story in {LANGUAGE}, a beginning, a continuation A and some subtle details.
Your task is to write a nuanced different continuation B of the story:

- The continuation must naturally **follow from the story so far**, not repeat or revise the story.
- Subtly but **accurately** weave all of the “follow-up” and “conflict” details provided into the main storyline with **minimal exposition**.
- Do **not** include, mention, explain, or describe the knowledge and premise in the provided details.
- Do **not** justify, rationalize, or explain the follow-up and conflict in the provided details.
- Focus on objective, observable descriptions of actions, settings, and dialogue.
- Do **NOT** add thoughts, feelings, or commentary.
- Do **NOT** mention events or outcomes that did **not** happen — focus on what is occurring in the scene.
- Do **NOT** use negation words (e.g., “not”, “no”, “never”, “didn’t”, “hasn’t”, “without”, “failed to”).
- Do **NOT** use speculative or ambiguous expressions (e.g., “seems”, “appears”, “perhaps”, “maybe”, “apparently”, “as if”).
- The continuation should be in **2–3 sentences**.

Example

Schema:

Scene: A quiet café in the afternoon.

Characters: Lisa, a journalist; Mark, her childhood friend.

Objects: Coffee cups, a notepad, a scarf hanging behind Lisa’s chair.

Situation description: Lisa leans across the table, her voice low as she asks Mark about the article. He listens, occasionally glancing at the entrance. The café hums softly around them.

Goal or activity: They are catching up and discussing a sensitive interview.

Visible tension or obstacle: Lisa is worried someone may overhear them.

Story so far:

Lisa lowered her voice, scribbling something in her notepad as Mark leaned in. The scarf behind her chair fluttered slightly as the door opened. He looked up, eyes scanning the new arrival. “Do you think they’re listening?” she whispered.

Continuation A:

Mark shook his head subtly, his eyes drifting past Lisa toward the window. A delivery man stepped inside, pausing to check the receipt in his hand.

Subtle details:

As they chatted, one of them quietly described the suspicious figure sneaking up behind himself.

Continuation B in English:

Mark nodded toward the hallway. “Someone just slipped behind me,” Lisa said, frowning.

Now you generate

Schema:

{schema}

Story so far:

{story}

Continuation A:

{silver}

Subtle details:

{details}

Continuation B in {LANGUAGE}: